

Data Summarization

Recap

- `select()`: subset and/or reorder columns
- `filter()`: remove rows
- `arrange()`: reorder rows
- `mutate()`: create new columns or modify them
- `select()` and `filter()` can be combined together
- remove a column: `select()` with `!` mark (`!col_name`)
- you can do sequential steps: especially using pipes `%>%`

▢ [Cheatsheet](#)

Another Cheatsheet

<https://raw.githubusercontent.com/rstudio/cheatsheets/main/data-transformation.pdf>

Data transformation with dplyr : : CHEAT SHEET



dplyr functions work with pipes and expect **tidy data**. In tidy data:



Each **variable** is in its own **column**



Each **observation**, or **case**, is in its own **row**



x %>% f(y) becomes **f(x, y)**

Summarise Cases

Apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).

summary function




summarise(.data, ...)
Compute table of summaries.
`summarise(mtcars, avg = mean(mpg))`



count(.data, ..., wt = NULL, sort = FALSE, name = NULL) Count number of rows in each group defined by the variables in ... Also **tally()**.
`count(mtcars, cyl)`

Group Cases

Use **group_by(.data, ..., .add = FALSE, .drop = TRUE)** to create a "grouped" copy of a table grouped by columns in ... dplyr functions will manipulate each "group" separately and combine the results.



`mtcars %>%
group_by(cyl) %>%
summarise(avg = mean(mpg))`


Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.




filter(.data, ..., .preserve = FALSE) Extract rows that meet logical criteria.
`filter(mtcars, mpg > 20)`



distinct(.data, ..., .keep_all = FALSE) Remove rows with duplicate values.
`distinct(mtcars, gear)`



slice(.data, ..., .preserve = FALSE) Select rows by position.
`slice(mtcars, 10:15)`



slice_sample(.data, ..., n, prop, weight_by = NULL, replace = FALSE) Randomly select rows. Use `n` to select a number of rows and `prop` to select a fraction of rows.
`slice_sample(mtcars, n = 5, replace = TRUE)`



slice_min(.data, order_by, ..., n, prop, with_ties = TRUE) and **slice_max()** Select rows with the lowest and highest values.
`slice_min(mtcars, mpg, prop = 0.25)`



slice_head(.data, ..., n, prop) and **slice_tail()** Select the first or last rows.
`slice_head(mtcars, n = 5)`

Logical and boolean operators to use with filter()

<code>==</code>	<code><</code>	<code><=</code>	<code>is.na()</code>	<code>%in%</code>	<code> </code>	<code>xor()</code>
<code>!=</code>	<code>></code>	<code>>=</code>	<code>!is.na()</code>	<code>!</code>	<code>&</code>	

See **?base::Logic** and **?Comparison** for help.

ARRANGE CASES

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.



pull(.data, var = -1, name = NULL, ...) Extract column values as a vector, by name or index.
`pull(mtcars, wt)`



select(.data, ...) Extract columns as a table.
`select(mtcars, mpg, wt)`



relocate(.data, ..., .before = NULL, .after = NULL) Move columns to new position.
`relocate(mtcars, mpg, cyl, .after = last_col())`

Use these helpers with select() and across()

e.g. `select(mtcars, mpg:cyl)`

contains(match)	num_range(prefix, range)	∴ e.g. <code>mpg:cyl</code>
ends_with(match)	all_of(x)/any_of(x, ..., vars)	∴ e.g. <code>-gear</code>
starts_with(match)	matches(match)	everything()

MANIPULATE MULTIPLE VARIABLES AT ONCE



across(.cols, funs, ..., .names = NULL) Summarise or mutate multiple columns in the same way.
`summarise(mtcars, across(everything(), mean))`



c_across(.cols) Compute across columns in row-wise data.
`transmute(rowwise(UKgas), total = sum(c_across(1:2)))`

MAKE NEW VARIABLES

Apply **vectorized functions** to columns. Vectorized functions take vectors as input and return vectors of the same length as output (see back)

Data Summarization

- Basic statistical summarization
 - `mean(x)`: takes the mean of x
 - `sd(x)`: takes the standard deviation of x
 - `median(x)`: takes the median of x
 - `quantile(x)`: displays sample quantiles of x. Default is min, IQR, max
 - `range(x)`: displays the range. Same as `c(min(x), max(x))`
 - `sum(x)`: sum of x
 - `max(x)`: maximum value in x
 - `min(x)`: minimum value in x
- **all have the `na.rm` = argument for missing data**

Statistical summarization

The vector getting summarized goes inside the parentheses:

```
x <- c(1, 5, 7, 4, 2, 8)
```

```
mean(x)
```

```
[1] 4.5
```

```
range(x)
```

```
[1] 1 8
```

```
sum(x)
```

```
[1] 27
```

Statistical summarization

Note that many of these functions have additional inputs regarding missing data, typically requiring the `na.rm` argument (“remove NAs”).

```
x <- c(1, 5, 7, 4, 2, 8, NA)
mean(x)
```

```
[1] NA
```

```
mean(x, na.rm = TRUE)
```

```
[1] 4.5
```

```
quantile(x)
```

```
Error in quantile.default(x): missing values and NaN's not allowed if 'na.rm' is FALSE
```

```
quantile(x, na.rm = TRUE)
```

0%	25%	50%	75%	100%
1.0	2.5	4.5	6.5	8.0

Statistical summarization

We will talk more about data types later, but you can only do summarization on numeric or logical types. Not characters.

```
x <- c(1, 5, 7, 4, 2, 8)
sum(x)
```

```
[1] 27
```

```
y <- c(TRUE, FALSE, FALSE, TRUE) # FALSE == 0 and TRUE == 1
sum(y)
```

```
[1] 2
```

```
z <- c("TRUE", "FALSE", "FALSE", "TRUE")
sum(z)
```

```
Error in sum(z): invalid 'type' (character) of argument
```

Some examples

We can use the `mtcars` built-in dataset. “The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973-74 models).”

The `head` command displays the first rows of an object:

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

The **dplyr** pipe `%>%` operator

A nice and readable way to chain together multiple R functions.

Changes `f(x, y)` to `x %>% f(y)`.

How mornings look like for most people:

```
me %>%  
  wake_up() %>%  
  get_out_of_bed() %>%  
  get_dressed() %>%  
  leave_house()
```

How my mornings look like most of the time:

```
leave_house(get_dressed(get_out_of_bed(wake_up(me))))
```

Statistical summarization the “tidy” way

```
mtcars %>% pull(hp) %>% mean() # alt: pull(mtcars, hp) %>% mean()
```

```
[1] 146.6875
```

```
mtcars %>% pull(wt) %>% median()
```

```
[1] 3.325
```

```
mtcars %>% pull(hp) %>% quantile()
```

```
   0%   25%   50%   75%  100%  
52.0  96.5 123.0 180.0 335.0
```

```
mtcars %>% pull(wt) %>% quantile(probs = 0.6)
```

```
   60%  
3.44
```

Behavior of `pull()` function

`pull()` converts a single data column into a [vector](#). This allows you to run summary functions on these data. Once you have “pulled” the data column out, you don’t have to name it again in any piped summary functions.

```
cars_wt <- mtcars %>% pull(wt)
class(cars_wt)
```

```
[1] "numeric"
```

```
cars_wt
```

```
 [1] 2.620 2.875 2.320 3.215 3.440 3.460 3.570 3.190 3.150 3.440 3.440 4.070
[13] 3.730 3.780 5.250 5.424 5.345 2.200 1.615 1.835 2.465 3.520 3.435 3.840
[25] 3.845 1.935 2.140 1.513 3.170 2.770 3.570 2.780
```

```
mtcars %>% pull(wt) %>% range(wt) # Incorrect
```

```
mtcars %>% pull(wt) %>% range() # Correct
```

```
[1] 1.513 5.424
```

GUT CHECK

What kind of object do we need to run summary operators like `mean()` ?

- A. A vector of numbers
- B. A vector of characters
- C. A dataset

Summarization on tibbles (data frames)

TB incidence

Let's read in a `tibble` of values from TB incidence.

"Tuberculosis incidence, all forms (per 100,000 population per year), for the period 1990-2007 across 208 countries/territories."

```
tb <- read_csv("https://jhudatascience.org/intro_to_r/data/tb.csv")
```

TB incidence

Check out the data:

```
head(tb)
```

```
# A tibble: 6 × 19
  TB incidence, all fo...1 `1990` `1991` `1992` `1993` `1994` `1995` `1996` `1997`
  <chr>          <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 Afghanistan    168   168   168   168   168   168   168   168
2 Albania         25    24    25    26    26    27    27    28
3 Algeria         38    38    39    40    41    42    43    44
4 American Samoa  21     7     2     9     9    11     0    12
5 Andorra         36    34    32    30    29    27    26    26
6 Angola         205   209   214   218   222   226   231   236
#   abbreviated name:
#   1`TB incidence, all forms (per 100 000 population per year)`
#   10 more variables: `1998` <dbl>, `1999` <dbl>, `2000` <dbl>, `2001` <dbl>,
#   `2002` <dbl>, `2003` <dbl>, `2004` <dbl>, `2005` <dbl>, `2006` <dbl>,
#   `2007` <dbl>
```

TB incidence

Check out the data:

```
str(tb)
```

```
spec_tbl_ [208 × 19] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ TB incidence, all forms (per 100 000 population per year): chr [1:208] "Afghanistan" "Albania"
 $ 1990 : num [1:208] 168 25 38 21 36 205 2
 $ 1991 : num [1:208] 168 24 38 7 34 209 24
 $ 1992 : num [1:208] 168 25 39 2 32 214 24
 $ 1993 : num [1:208] 168 26 40 9 30 218 24
 $ 1994 : num [1:208] 168 26 41 9 29 222 23
 $ 1995 : num [1:208] 168 27 42 11 27 226 2
 $ 1996 : num [1:208] 168 27 43 0 26 231 23
 $ 1997 : num [1:208] 168 28 44 12 26 236 2
 $ 1998 : num [1:208] 168 28 46 6 25 240 23
 $ 1999 : num [1:208] 168 27 47 8 23 245 23
 $ 2000 : num [1:208] 168 25 48 6 22 250 23
 $ 2001 : num [1:208] 168 23 49 6 21 255 22
 $ 2002 : num [1:208] 168 23 50 4 21 260 22
 $ 2003 : num [1:208] 168 22 51 5 20 265 22
 $ 2004 : num [1:208] 168 21 53 9 20 270 22
 $ 2005 : num [1:208] 168 20 54 10 19 276 2
 $ 2006 : num [1:208] 168 18 55 7 19 281 22
 $ 2007 : num [1:208] 168 17 57 5 19 287 22
 - attr(*, "spec")=
 .. cols(
 ..   `TB incidence, all forms (per 100 000 population per year)` = col_character(),
 ..   `1990` = col_double(),
 ..   `1991` = col_double(),
```


Indicator of TB

Before we go further, let's rename the first column using the `rename()` function in `dplyr`.

In this case, we have to use the backticks (``) because there are spaces and funky characters in the name.

```
tb <- tb %>%  
  rename(country = `TB incidence, all forms (per 100 000 population per year)`)
```

Indicator of TB

`colnames()` will show us the column names and show that country is renamed:

```
colnames(tb)
```

```
[1] "country" "1990"    "1991"    "1992"    "1993"    "1994"    "1995"  
[8] "1996"    "1997"    "1998"    "1999"    "2000"    "2001"    "2002"  
[15] "2003"    "2004"    "2005"    "2006"    "2007"
```

Summarize the data: **dplyr** `summarize()` function

`summarize` creates a summary table of a column you're interested in.

Can run multiple summary statistics at once (unlike `pull()` which can only do a single calculation on one column).

You can also do more elaborate summaries across different groups of data using `group_by()`. More on this later!

General format - Not the code!

```
{data to use} %>%  
  summarize({summary column name} = {function(source column)},  
            {summary column name} = {function(source column)})
```

Summarize the data: `dplyr` `summarize()` function

`summarize` creates a summary table of a column you're interested in.

General format - Not the code!

```
{data to use} %>%  
  summarize({summary column name} = {function(source column)})
```

tb %>%

```
  summarize(mean_1991 = mean(`1991`)) # Note the backticks, this is a column name!
```

A tibble: 1 × 1

```
  mean_1991  
    <dbl>  
1         NA
```

tb %>%

```
  summarize(mean_1991 = mean(`1991`, na.rm = TRUE))
```

A tibble: 1 × 1

```
  mean_1991  
    <dbl>  
1      108.
```

Summarize the data: `dplyr summarize()` function

`summarize()` can do multiple operations at once. Just separate by a comma.

```
tb %>%  
  summarize(mean_1991 = mean(`1991`, na.rm = TRUE),  
            median_1991 = median(`1991`, na.rm = TRUE),  
            median(`2000`, na.rm = TRUE))  
  
# A tibble: 1 × 3  
  mean_1991 median_1991 `median(\`2000\`, na.rm = TRUE)`  
    <dbl>      <dbl>                <dbl>  
1    108.         58                60
```

Notice how when we forget to provide a new name, output is still provided, but the column name is messy.

Summarize the data: `dplyr summarize()` function

This looks better.

```
tb %>%  
  summarize(mean_1991 = mean(`1991`, na.rm = TRUE),  
            median_1991 = median(`1991`, na.rm = TRUE),  
            median_2000 = median(`2000`, na.rm = TRUE))
```

```
# A tibble: 1 × 3  
  mean_1991 median_1991 median_2000  
    <dbl>      <dbl>      <dbl>  
1    108.         58         60
```

Summarize the data: **dplyr** `summarize()` function

Note that `summarize()` creates a separate tibble from the original data, so you don't want to overwrite your original data if you decide to save the summary.

If you want to save a summary statistic in the original data, use `mutate()` instead to create a new column for the summary statistic.

summary() Function

Using `summary()` can give you rough snapshots of each numeric column (character columns are skipped):

```
summary(tb)
```

country	1990	1991	1992		
Length:208	Min. : 0.0	Min. : 4.0	Min. : 2.0		
Class :character	1st Qu.: 27.5	1st Qu.: 27.0	1st Qu.: 27.0		
Mode :character	Median : 60.0	Median : 58.0	Median : 56.0		
	Mean :105.6	Mean :107.7	Mean :108.3		
	3rd Qu.:165.0	3rd Qu.:171.0	3rd Qu.:171.5		
	Max. :585.0	Max. :594.0	Max. :606.0		
	NA's :1	NA's :1	NA's :1		
1993	1994	1995	1996	1997	
Min. : 4.0	Min. : 0	Min. : 3.0	Min. : 0.0	Min. : 0.0	
1st Qu.: 27.5	1st Qu.: 26	1st Qu.: 26.5	1st Qu.: 25.5	1st Qu.: 24.5	
Median : 56.0	Median : 57	Median : 58.0	Median : 60.0	Median : 64.0	
Mean :110.3	Mean :112	Mean :114.2	Mean :115.4	Mean :118.9	
3rd Qu.:171.0	3rd Qu.:174	3rd Qu.:177.5	3rd Qu.:179.0	3rd Qu.:181.0	
Max. :618.0	Max. :630	Max. :642.0	Max. :655.0	Max. :668.0	
NA's :1	NA's :1	NA's :1	NA's :1	NA's :1	
1998	1999	2000	2001		
Min. : 0.0	Min. : 0.0	Min. : 0.0	Min. : 0.0		
1st Qu.: 23.5	1st Qu.: 22.5	1st Qu.: 21.5	1st Qu.: 19.0		
Median : 63.0	Median : 66.0	Median : 60.0	Median : 59.0		
Mean :121.5	Mean :125.0	Mean :127.8	Mean :130.7		
3rd Qu.:188.5	3rd Qu.:192.5	3rd Qu.:191.0	3rd Qu.:189.5		
Max. :681.0	Max. :695.0	Max. :801.0	Max. :916.0		
NA's :1	NA's :1	NA's :1	NA's :1		

Summary & Lab Part 1

- `pull()` creates a *vector*
- don't forget the `na.rm = TRUE` argument!
- `summary(x)`: quantile information
- `summarize`: creates a summary table of columns of interest
- summary stats (`mean()`) work with vectors or with `summarize()`

▢ [Class Website](#)

▢ [Lab](#)

▢ [Day 4 Cheatsheet](#)

Colorado heat-related ER visits

Here we will be using data on the annual rates of medical visits associated with heat-related illness in Colorado:

http://jhudatascience.org/intro_to_r/data/CO_ER_heat_visits_2011_2024.csv

- Check out info about the data at: <https://coepht.colorado.gov/heat-data>

```
CO_heat <- read_csv("http://jhudatascience.org/intro_to_r/data/CO_ER_heat_visits_2011_2024.csv")
head(CO_heat)
```

```
# A tibble: 6 × 11
```

	COUNTY	RATE	L95CL	U95CL	VISITS	FLAG	YEAR	GENDER	MEASURE	HEALTHOUTCOMEID
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<chr>
1	Statewide	5.65	4.70	6.60	140	2	2011	Female	Age ad...	Heat-related
2	Statewide	7.39	6.30	8.48	183	2	2011	Male	Age ad...	Heat-related
3	Statewide	6.52	5.80	7.24	323	2	2011	Both g...	Age ad...	Heat-related
4	Statewide	5.65	4.72	6.58	146	2	2012	Female	Age ad...	Heat-related
5	Statewide	7.57	6.48	8.65	193	2	2012	Male	Age ad...	Heat-related
6	Statewide	6.59	5.88	7.30	339	2	2012	Both g...	Age ad...	Heat-related

```
#   1 more variable: cofips <chr>
```

distinct() values

`distinct(x)` will return the unique elements of column x.

```
CO_heat %>%  
  distinct(COUNTY)
```

```
# A tibble: 65 × 1
```

```
  COUNTY
```

```
  <chr>
```

```
1 Statewide
```

```
2 Adams
```

```
3 Alamosa
```

```
4 Arapahoe
```

```
5 Archuleta
```

```
6 Baca
```

```
7 Bent
```

```
8 Boulder
```

```
9 Broomfield
```

```
10 Chaffee
```

```
#   55 more rows
```

How many `distinct()` values?

`n_distinct()` tells you the number of unique elements. It needs a vector so you *must pull the column first!*

```
CO_heat %>%  
  pull(COUNTY) %>%  
  n_distinct()
```

```
[1] 65
```

Use `count()` to return row count per category.

Use `count` to return a frequency table of unique elements of a data.frame.

```
CO_heat %>% count(COUNTY)
```

```
# A tibble: 65 × 2
  COUNTY      n
  <chr>    <int>
1 Adams      42
2 Alamosa     42
3 Arapahoe    42
4 Archuleta   42
5 Baca        42
6 Bent        42
7 Boulder     42
8 Broomfield  42
9 Chaffee     42
10 Cheyenne   42
#   55 more rows
```

Multiple columns listed further subdivides the `count()`

```
CO_heat %>% count(COUNTY, GENDER)
```

```
# A tibble: 195 × 3
```

	COUNTY <chr>	GENDER <chr>	n <int>
1	Adams	Both genders	14
2	Adams	Female	14
3	Adams	Male	14
4	Alamosa	Both genders	14
5	Alamosa	Female	14
6	Alamosa	Male	14
7	Arapahoe	Both genders	14
8	Arapahoe	Female	14
9	Arapahoe	Male	14
10	Archuleta	Both genders	14

```
#   185 more rows
```

Note: `count()` includes NAs

GUT CHECK

The `count()` function can help us tally:

- A. Sample size
- B. Rows per each category
- C. How many categories

Grouping

Goal

We want to find the average rate for ER visits for heat-related illness in the dataset.

How do we do this?

Perform operations By groups: dplyr

`group_by` allows you group the data set by variables/columns you specify:

```
# Regular data
```

```
CO_heat
```

```
# A tibble: 2,730 × 11
```

	COUNTY	RATE	L95CL	U95CL	VISITS	FLAG	YEAR	GENDER	MEASURE	HEALTHOUTCOMEID
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<chr>
1	Statewide	5.65	4.70	6.60	140	2	2011	Female	Age ad...	Heat-related
2	Statewide	7.39	6.30	8.48	183	2	2011	Male	Age ad...	Heat-related
3	Statewide	6.52	5.80	7.24	323	2	2011	Both ...	Age ad...	Heat-related
4	Statewide	5.65	4.72	6.58	146	2	2012	Female	Age ad...	Heat-related
5	Statewide	7.57	6.48	8.65	193	2	2012	Male	Age ad...	Heat-related
6	Statewide	6.59	5.88	7.30	339	2	2012	Both ...	Age ad...	Heat-related
7	Statewide	4.96	4.08	5.84	124	2	2013	Female	Age ad...	Heat-related
8	Statewide	6.70	5.70	7.70	178	2	2013	Male	Age ad...	Heat-related
9	Statewide	5.83	5.16	6.49	302	2	2013	Both ...	Age ad...	Heat-related
10	Statewide	3.54	2.81	4.27	92	2	2014	Female	Age ad...	Heat-related

```
# 2,720 more rows
```

```
# 1 more variable: cofips <chr>
```

Perform operations by groups: dplyr

`group_by` allows you group the data set by variables/columns you specify:

```
CO_heat_grouped <- CO_heat %>% group_by(COUNTY)
CO_heat_grouped
```

```
# A tibble: 2,730 × 11
```

```
# Groups:   COUNTY [65]
```

	COUNTY	RATE	L95CL	U95CL	VISITS	FLAG	YEAR	GENDER	MEASURE	HEALTHOUTCOMEID
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<chr>
1	Statewide	5.65	4.70	6.60	140	2	2011	Female	Age ad...	Heat-related
2	Statewide	7.39	6.30	8.48	183	2	2011	Male	Age ad...	Heat-related
3	Statewide	6.52	5.80	7.24	323	2	2011	Both ...	Age ad...	Heat-related
4	Statewide	5.65	4.72	6.58	146	2	2012	Female	Age ad...	Heat-related
5	Statewide	7.57	6.48	8.65	193	2	2012	Male	Age ad...	Heat-related
6	Statewide	6.59	5.88	7.30	339	2	2012	Both ...	Age ad...	Heat-related
7	Statewide	4.96	4.08	5.84	124	2	2013	Female	Age ad...	Heat-related
8	Statewide	6.70	5.70	7.70	178	2	2013	Male	Age ad...	Heat-related
9	Statewide	5.83	5.16	6.49	302	2	2013	Both ...	Age ad...	Heat-related
10	Statewide	3.54	2.81	4.27	92	2	2014	Female	Age ad...	Heat-related

```
#   2,720 more rows
```

```
#   1 more variable: cofips <chr>
```

Summarize the grouped data

It's grouped! Grouping doesn't change the data in any way, but how **functions operate on it**. Now we can summarize **RATE** (age-adjusted heat-related ER visits per 100,000 people) by group:

```
CO_heat_grouped %>% summarize(avg_rate = mean(RATE, na.rm = TRUE))
```

```
# A tibble: 65 × 2
  COUNTY      avg_rate
  <chr>      <dbl>
1 Adams      6.85
2 Alamosa     0
3 Arapahoe   5.20
4 Archuleta   0
5 Baca       0
6 Bent       0
7 Boulder    5.46
8 Broomfield  0
9 Chaffee    0
10 Cheyenne   0
#   55 more rows
```

Do it in one step: use %>% to string these together!

Pipe CO_heat into group_by, then pipe that into summarize:

```
CO_heat %>%  
  group_by(COUNTY) %>%  
  summarize(avg_rate = mean(RATE, na.rm = TRUE),  
            max_rate = max(RATE, na.rm = TRUE))
```

```
# A tibble: 65 × 3
```

	COUNTY <chr>	avg_rate <dbl>	max_rate <dbl>
1	Adams	6.85	11.7
2	Alamosa	0	0
3	Arapahoe	5.20	8.44
4	Archuleta	0	0
5	Baca	0	0
6	Bent	0	0
7	Boulder	5.46	8.96
8	Broomfield	0	0
9	Chaffee	0	0
10	Cheyenne	0	0

```
#   55 more rows
```

Group by as many variables as you want

group_by COUNTY and GENDER:

```
CO_heat %>%  
  group_by(COUNTY, GENDER) %>%  
  summarize(avg_rate = mean(RATE, na.rm = TRUE),  
            max_rate = max(RATE, na.rm = TRUE))
```

Warning: There were 4 warnings in `summarize()`.

The first warning was:

- ▮ In argument: `max\_rate = max(RATE, na.rm = TRUE)`.
- ▮ In group 20: `COUNTY = "Boulder"` `GENDER = "Female"`.

Caused by warning in `max()`:

! no non-missing arguments to max; returning -Inf

- ▮ Run `dplyr::last\_dplyr\_warnings()` to see the 3 remaining warnings.

A tibble: 195 × 4

Groups: COUNTY [65]

	COUNTY	GENDER	avg_rate	max_rate
	<chr>	<chr>	<dbl>	<dbl>
1	Adams	Both genders	6.35	8.96
2	Adams	Female	6.37	7.59
3	Adams	Male	7.83	11.7
4	Alamosa	Both genders	0	0
5	Alamosa	Female	0	0
6	Alamosa	Male	0	0
7	Arapahoe	Both genders	4.93	7.16
8	Arapahoe	Female	4.77	5.86
9	Arapahoe	Male	5.94	8.44

Only the last `group_by` is recognized...

You can overwrite the first `group_by` with a new one.

```
CO_heat %>%
  group_by(COUNTY, GENDER) %>%
  group_by(GENDER)

# A tibble: 2,730 × 11
# Groups:   GENDER [3]
  COUNTY     RATE L95CL U95CL VISITS  FLAG  YEAR GENDER MEASURE HEALTHOUTCOMEID
  <chr>     <dbl> <dbl> <dbl>   <dbl> <dbl> <dbl> <chr>  <chr>    <chr>
1 Statewide  5.65  4.70  6.60    140     2  2011 Female Age ad... Heat-related
2 Statewide  7.39  6.30  8.48    183     2  2011 Male   Age ad... Heat-related
3 Statewide  6.52  5.80  7.24    323     2  2011 Both ... Age ad... Heat-related
4 Statewide  5.65  4.72  6.58    146     2  2012 Female Age ad... Heat-related
5 Statewide  7.57  6.48  8.65    193     2  2012 Male   Age ad... Heat-related
6 Statewide  6.59  5.88  7.30    339     2  2012 Both ... Age ad... Heat-related
7 Statewide  4.96  4.08  5.84    124     2  2013 Female Age ad... Heat-related
8 Statewide  6.70  5.70  7.70    178     2  2013 Male   Age ad... Heat-related
9 Statewide  5.83  5.16  6.49    302     2  2013 Both ... Age ad... Heat-related
10 Statewide 3.54  2.81  4.27     92     2  2014 Female Age ad... Heat-related
#   2,720 more rows
#   1 more variable: cofips <chr>
```

Ungroup the data

The ungroup function will allow you to clear the groups from the data.

```
CO_heat_grouped
```

```
# A tibble: 2,730 × 11
# Groups:   COUNTY [65]
  COUNTY    RATE L95CL U95CL VISITS FLAG  YEAR GENDER MEASURE HEALTHOUTCOMEID
  <chr>    <dbl> <dbl> <dbl>   <dbl> <dbl> <dbl> <chr>  <chr>    <chr>
1 Statewide  5.65  4.70  6.60    140     2  2011 Female Age ad... Heat-related
2 Statewide  7.39  6.30  8.48    183     2  2011 Male   Age ad... Heat-related
3 Statewide  6.52  5.80  7.24    323     2  2011 Both ... Age ad... Heat-related
4 Statewide  5.65  4.72  6.58    146     2  2012 Female Age ad... Heat-related
5 Statewide  7.57  6.48  8.65    193     2  2012 Male   Age ad... Heat-related
6 Statewide  6.59  5.88  7.30    339     2  2012 Both ... Age ad... Heat-related
7 Statewide  4.96  4.08  5.84    124     2  2013 Female Age ad... Heat-related
8 Statewide  6.70  5.70  7.70    178     2  2013 Male   Age ad... Heat-related
9 Statewide  5.83  5.16  6.49    302     2  2013 Both ... Age ad... Heat-related
10 Statewide 3.54  2.81  4.27     92     2  2014 Female Age ad... Heat-related
#   2,720 more rows
#   1 more variable: cofips <chr>
```

```
ungroup(CO_heat_grouped)
```

```
# A tibble: 2,730 × 11
  COUNTY    RATE L95CL U95CL VISITS FLAG  YEAR GENDER MEASURE HEALTHOUTCOMEID
  <chr>    <dbl> <dbl> <dbl>   <dbl> <dbl> <dbl> <chr>  <chr>    <chr>
1 Statewide  5.65  4.70  6.60    140     2  2011 Female Age ad... Heat-related
2 Statewide  7.39  6.30  8.48    183     2  2011 Male   Age ad... Heat-related
3 Statewide  6.52  5.80  7.24    323     2  2011 Both ... Age ad... Heat-related
```


group_by with mutate - just add data

We can also use `mutate` to calculate the mean value for each year and add it as a column:

```
CO_heat %>%  
  group_by(YEAR, GENDER) %>%  
  mutate(year_avg = mean(RATE, na.rm = TRUE)) %>%  
  select(COUNTY, RATE, year_avg)
```

```
# A tibble: 2,730 × 5
```

```
# Groups:   YEAR, GENDER [42]
```

	YEAR	GENDER	COUNTY	RATE	year_avg
	<dbl>	<chr>	<chr>	<dbl>	<dbl>
1	2011	Female	Statewide	5.65	0.950
2	2011	Male	Statewide	7.39	1.38
3	2011	Both genders	Statewide	6.52	1.59
4	2012	Female	Statewide	5.65	0.891
5	2012	Male	Statewide	7.57	1.71
6	2012	Both genders	Statewide	6.59	1.86
7	2013	Female	Statewide	4.96	0.795
8	2013	Male	Statewide	6.70	1.65
9	2013	Both genders	Statewide	5.83	1.92
10	2014	Female	Statewide	3.54	0.508

```
#   2,720 more rows
```

Counting

There are other functions, such as `n()` count the number of observations (NAs included).

```
CO_heat %>%  
  group_by(YEAR) %>%  
  summarize(n = n(),  
            mean = mean(RATE, na.rm = TRUE))
```

```
# A tibble: 14 × 3  
  YEAR      n mean  
  <dbl> <int> <dbl>  
1  2011    195 1.29  
2  2012    195 1.50  
3  2013    195 1.45  
4  2014    195 0.967  
5  2015    195 1.54  
6  2016    195 3.19  
7  2017    195 1.55  
8  2018    195 2.84  
9  2019    195 3.34  
10 2020    195 1.26  
11 2021    195 1.68  
12 2022    195 2.45  
13 2023    195 2.60  
14 2024    195 2.62
```

Counting

`count()` and `n()` can give very similar information.

```
CO_heat %>% count(YEAR) %>% head(n = 3)
```

```
# A tibble: 3 × 2
```

	YEAR	n
	<dbl>	<int>
1	2011	195
2	2012	195
3	2013	195

```
CO_heat %>% group_by(YEAR) %>% summarize(n = n()) %>% head(n = 3) # n() typically used with summarize
```

```
# A tibble: 3 × 2
```

	YEAR	n
	<dbl>	<int>
1	2011	195
2	2012	195
3	2013	195

A few miscellaneous topics

Base R functions you might see: **length** and **unique**

These functions require a column as a vector using `pull()`.

```
CO_heat_loc <- CO_heat %>% pull(COUNTY) # pull() to make a vector
CO_heat_loc %>% unique() # similar to distinct()
```

```
[1] "Statewide"  "Adams"      "Alamosa"    "Arapahoe"   "Archuleta"
[6] "Baca"       "Bent"       "Boulder"    "Broomfield" "Chaffee"
[11] "Cheyenne"   "Clear Creek" "Conejos"    "Costilla"   "Crowley"
[16] "Custer"     "Delta"      "Denver"     "Dolores"    "Douglas"
[21] "Eagle"      "Elbert"     "El Paso"    "Fremont"    "Garfield"
[26] "Gilpin"     "Grand"      "Gunnison"   "Hinsdale"   "Huerfano"
[31] "Jackson"    "Jefferson"  "Kiowa"      "Kit Carson" "Lake"
[36] "La Plata"   "Larimer"    "Las Animas" "Lincoln"    "Logan"
[41] "Mesa"       "Mineral"    "Moffat"     "Montezuma"  "Montrose"
[46] "Morgan"     "Otero"      "Ouray"      "Park"       "Phillips"
[51] "Pitkin"     "Prowers"    "Pueblo"     "Rio Blanco" "Rio Grande"
[56] "Routt"      "Saguache"   "San Juan"   "San Miguel" "Sedgwick"
[61] "Summit"     "Teller"     "Washington" "Weld"       "Yuma"
```

Base R functions you might see: **length** and **unique**

These functions require a column as a vector using `pull()`.

```
CO_heat_loc %>% unique() %>% length() # similar to n_distinct()
```

```
[1] 65
```

summary() vs. summarize()

- `summary()` (base R) gives statistics table on a dataset.
- `summarize()` (dplyr) creates a more customized summary tibble/dataframe.

Functions you might also see

- `sum(!is.na())`: # of non-NAs in the data
- `first()`: first value in the data
- `last()`: last value in the data
- `range()`: minimum and maximum of the data
- `IQR()`: interquartile range of the data

Summary

- `count(x)`: what unique values do you have?
 - `distinct()`: what are the distinct values?
 - `n_distinct()` with `pull()`: how many distinct values?
- `group_by()`: changes subsequent functions (remove with `ungroup()`)
 - combine with `summarize()` to get statistics per group
 - combine with `mutate()` to add column
- `summarize()` with `n()` gives the count (NAs included)

Resources & Lab Part 2

- ▮ [Class Website](#)
- ▮ [Lab](#)
- ▮ [Day 4 Cheatsheet](#)
- ▮ [Posit's data transformation Cheatsheet](#)



Image by [Gerd Altmann](#) from [Pixabay](#)

**Extra Slides: More advanced
summarization**

Data Summarization on data frames

- Statistical summarization across the data frame
 - `rowMeans(x)`: takes the means of each row of x
 - `colMeans(x)`: takes the means of each column of x
 - `rowSums(x)`: takes the sum of each row of x
 - `colSums(x)`: takes the sum of each column of x

rowMeans() example

Get means for each row.

Let's see what the mean TB incidence is across years each row (country):

```
tb %>%  
  select(starts_with("year")) %>%  
  rowMeans(na.rm = TRUE) %>%  
  head(n = 5)
```

```
[1] NaN NaN NaN NaN NaN
```

```
tb %>%  
  group_by(country) %>%  
  summarize(mean = rowMeans(across(starts_with("year")), na.rm = TRUE)) %>%  
  head(n = 5)
```

```
# A tibble: 5 × 2  
  country      mean  
  <chr>      <dbl>  
1 Afghanistan  NaN  
2 Albania      NaN  
3 Algeria      NaN  
4 American Samoa NaN  
5 Andorra      NaN
```

colMeans() example

Get means for each column.

Let's see what the mean is across each column (year):

```
tb %>%  
  select(starts_with("year")) %>%  
  colMeans(na.rm = TRUE) %>%  
  head(n = 5)  
  
numeric(0)  
  
tb %>%  
  summarize(across(starts_with("year"), ~mean(.x, na.rm = TRUE)))  
  
# A tibble: 1 × 0
```

* New! * Many dplyr functions now have a **.by=** argument

Pipe `CO_heat` into `group_by`, then pipe that into `summarize`:

```
CO_heat %>%  
  group_by(COUNTY) %>%  
  summarize(avg_rate = mean(RATE, na.rm = TRUE),  
            max_rate = max(RATE, na.rm = TRUE))
```

is the same as..

```
CO_heat %>%  
  summarize(avg_rate = mean(RATE, na.rm = TRUE),  
            max_rate = max(RATE, na.rm = TRUE),  
            .by = COUNTY)
```